Laser Cutting Intelligence

Surface Roughness Prediction via Ridge Regression

"When your process variables move together, your model needs a steadying hand — that's what Ridge does."

0.722	0.394	4.422	7.6
R^2 Score	RMSE (µm)	Optimal Alpha	Max VIF (Ridge flag)
Ridge L2 Penalty	RidgeCV 5-fold	Multicollinearity	StandardScaler

Precision Fabrication · CNC Laser Cutting · Process Engineering · Surface Quality

Luis Lozano | Operational Excellence Manager · MBB · Machine Learning | GitHub: LozanoLsa

Turning Operations into Predictive Systems

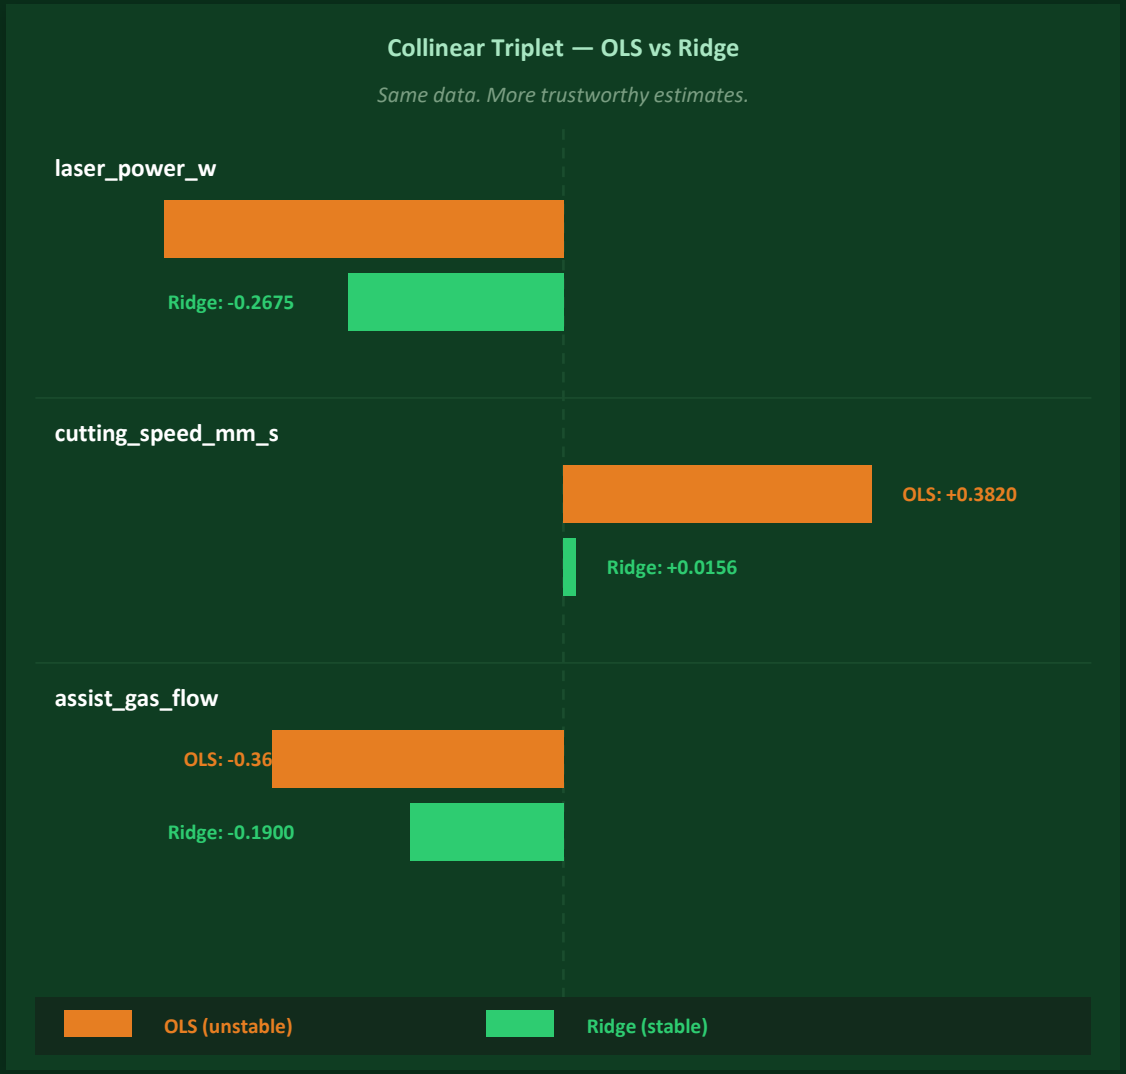

OLS sees correlated variables as a numerical problem. Ridge sees them as a structure to accommodate.

The Real Problem

The Multicollinearity Trap

Power, speed, and gas pressure don't move independently — operators increase them in concert because that is correct practice. This r = 0.907 correlation creates a near-singular design matrix that makes OLS coefficients unreliable. The model might suggest reducing power when the real issue is speed — and the engineer follows bad guidance.

Recipe Instability Without Regularisation

Without Ridge, a small change in the training sample — a new batch of cuts — can flip a coefficient's sign or double its magnitude. A model trained in winter that assigns a specific power coefficient cannot be trusted in summer when the training distribution shifts slightly. OLS has no stabilising mechanism for correlated inputs.

47.8% of Cuts Outside Ra Spec

Across 1,500 cut records, nearly half exceed the Ra ≤ 3.2 µm structural quality threshold. Carbon steel averages 3.48 µm mean Ra — already above spec. Without a model that can reliably guide recipe corrections, operators are adjusting by feel in a process where the key levers are mathematically entangled.

Insight: Ridge doesn't eliminate the correlation — it acknowledges it and distributes the coefficient effect proportionally across all correlated variables.

A model that gives wrong directions is more dangerous than no model at all.

From Unstable to Stable

BEFORE — OLS on Correlated Process Data

01 Fit OLS to 8 Features

Power, speed, and gas are all included. OLS attempts to partition their shared variance.

02 Near-Singular Design Matrix

r = 0.907 (power vs speed). OLS cannot stably invert the matrix. Coefficients amplify noise.

03 Inflated, Unstable Coefficients

OLS assigns laser_power −0.495 and cutting_speed +0.382 — both inflated vs true partial effect. Signs may flip with a new training batch.

04 Guidance Misleads the Engineer

"Reduce power" when the actual driver is speed. The model is accurate on average but wrong at the coefficient level.

05 No Convergence Across Batches

Quarterly model refit produces different coefficient signs. Process team loses trust. Model is abandoned.

VS

AFTER — Ridge on the Same Data

01 Same 8 Features, StandardScaler Applied

L2 penalty requires equitable scaling — laser_power (σ≈720) and focal_offset (σ≈0.6) on the same footing.

02 L2 Penalty Stabilises the Matrix

Ridge adds α·I to the design matrix before inversion — resolving near-singularity without discarding any variable.

03 Smooth, Conservative Coefficients

Ridge assigns laser_power −0.268, cutting_speed +0.016 — smaller, stable, physically plausible. Both keep their expected signs.

04 Guidance the Engineer Can Trust

Every coefficient direction is physically grounded. More power reduces Ra. Faster speed increases Ra. Always, stably.

05 Stable Across Quarterly Refits

RidgeCV re-selects alpha. Coefficients shift modestly. Process team can compare refits quarter-over-quarter.

Ridge doesn't solve the correlation — it makes the model behave sensibly in its presence.

The Data

Where Each Variable Lives in the Laser Cutting Cell

Variable	Source System	Operational Context
laser_power_w, cutting_speed_mm_s	CNC Laser Controller / G-Code Log	Recipe parameters from the cutting program file — logged at job start
assist_gas_flow_l_min, oxygen_pct	Gas Console / Flow Controller	Assist gas parameters set at job setup or material changeover
focal_offset_mm	CNC Laser Controller	Focal offset from cutting head config — set per recipe
material_thickness_mm	ERP / Production Order	Nominal sheet thickness from the job traveller or material spec
material_type	ERP / Material Master	Binary flag (0=steel, 1=aluminium) from job setup record
shop_temp_c	Environmental Sensor	Bay-level ambient temperature at cut time — data historian
surface_roughness_ra_um	Post-Cut CMM / Profilometer	TARGET — Ra measured at the cut edge after job completion

Joining this dataset requires connecting the CNC job log, the ERP material spec, the environmental data historian, and the CMM quality station on a common job ID and timestamp. The cut parameter log and the Ra measurement are rarely in the same system.

Dataset Profile

Records	1,500
Features	8
Train	1,200
Test	300
Ra Spec	≤ 3.2 µm
In-Spec	52.2%
Steel Ra avg	3.48 µm
Al Ra avg	2.86 µm
Max VIF	7.6 (power)

VIF > 5 is not a data quality warning — it is the quantified justification for Ridge over OLS.

Why Ridge?

"Ridge keeps all features in play — appropriate when every parameter has a physical justification for inclusion."

01 L2 Penalty Stabilises Without Zeroing

Unlike Lasso's L1 norm, the L2 norm never reaches exactly zero — it asymptotes toward it. This means Ridge keeps all 8 laser parameters and assigns each a stable partial effect. Appropriate when every variable has a physical reason to be in the model and should not be eliminated.

02 VIF > 5 Is the Ridge Diagnostic

Variance Inflation Factor of 7.6 for laser_power means OLS cannot reliably estimate its coefficient independently of speed and gas. Ridge adds $\alpha \cdot I$ to the design matrix before inversion — mathematically equivalent to solving the near-singularity problem. The VIF is not a warning; it is the design input.

03 RidgeCV Makes Alpha Selection Principled

80 candidates on a log scale [0.001, 1000], 5-fold cross-validation, R^2 scoring. Selected alpha = 4.422. This is the L2 penalty that maximises generalisation on held-out data — not the one that looks best on training set. The cross-validation prevents over-shrinkage.

04 Equal R^2 with OLS Is the Correct Outcome

Ridge R^2 = 0.722. OLS R^2 = 0.722. Identical. This is not a tie — it is proof that Ridge did not sacrifice predictive accuracy for stability. On well-specified data, Ridge should match OLS accuracy while producing more trustworthy coefficients. That is the deal.

Lasso zeroes the irrelevant. Ridge steadies the correlated. Choosing wrong doesn't break the model — it breaks its interpretability.

Model Results

0.722	0.394	0.305	4.422
R^2	RMSE (µm)	MAE (µm)	Optimal Alpha
72.2% variance explained	*12.3% of spec limit*	*Adequate for recipe scheduling*	*RidgeCV 5-fold selected*

Ridge Coefficients — Standardised, Stable Under Collinearity

material_thickness_mm	+0.385 /σ	Dominant — more material = energy gap → roughness
material_type	−0.304 /σ	Aluminium's conductivity → consistently lower Ra
laser_power_w	−0.268 /σ	More power → cleaner melt ejection → smoother edge
assist_gas_flow_l_min	−0.190 /σ	Gas assists melt removal — stable, not inflated
oxygen_pct	+0.052 /σ	Higher O_2 mildly increases edge oxidation
shop_temp_c	+0.025 /σ	Ambient variation — small consistent effect
cutting_speed_mm_s	+0.016 /σ	Faster travel → less energy density → slight Ra rise
focal_offset_mm	+0.006 /σ	Defocus reduces energy density — smallest effect

★ Green accent = collinear triplet (power · speed · gas). Ridge keeps all three stable.
Ridge matching OLS R^2 is not a tie — it is proof that stability costs nothing in predictive power.

Model Validation

Ridge R^2	0.722	*72.2% variance explained*
OLS R^2	0.722	*Identical — Ridge wins on stability*
Alpha CV	4.422	*RidgeCV 5-fold selected*
Max VIF	7.6	*Power: Ridge justified*
RMSE/Spec	12.3%	*0.394/3.2 µm*

From Insight to Action

01 — Process Engineer

Use the Model as a Pre-Cut Recipe Validator

Enter the 8 parameters for a proposed cutting recipe and get a predicted Ra before committing to the cut. If the model predicts Ra > 3.2 µm, adjust power (up) or speed (down) before the job starts — not after the CMM reports a defect.

IMPACT

Eliminates reactive deburring on out-of-spec cuts

02 — Quality Engineer

Use the 2D Response Surface as a Process Window Map

The Power × Speed contour plot for each material/thickness combination shows which recipes meet spec. Print it as a quick-reference card for operators — a visual specification limit that replaces tribal knowledge with quantified boundaries.

IMPACT

Standardises recipe selection across shifts and operators

03 — Manufacturing Engineer

Quantify the Thickness–Roughness Trade-off

Material thickness is the dominant driver (+0.385 per σ). The model makes the Ra cost of cutting thicker material explicit: going from 3 mm to 12 mm steel requires a recipe correction of approximately +1.3W power and −17 mm/s speed to stay within spec.

IMPACT

Physics-based recipe scaling for new material thicknesses

04 — OPEX / Lean Team

Target Coefficient Instability as a Process Risk

A model whose power coefficient flips sign between quarterly refits signals a process that has drifted — not a data problem. Use the regularisation path as a process health dashboard: stable coefficients across refits = stable process. Flip = investigate.

IMPACT

Ridge refit as a leading indicator of process drift

05 — ML Practitioner

Benchmark Ridge vs ElasticNet for the Next Iteration

Ridge keeps all features. Lasso zeroes some. ElasticNet (L1+L2) does both — and may perform better when the correlated triplet truly shares predictive content. Running all three on the same dataset and comparing their coefficient trajectories is the logical next experiment.

IMPACT

Maps the full regularisation landscape for this process

Stable coefficients are not a luxury — they are the minimum requirement for a model that guides process decisions.

What-If Simulator

Same spec limit, same material — different power and speed recipes. The model predicts what the profilometer will find.

Scenario A — 3 mm Aluminium · Optimised Recipe

Material	**Aluminium · 3 mm**
Laser Power	**3,200 W**
Cut Speed	**20 mm/s**
Gas Flow	**10 L/min**
O_2 Content	**8%**

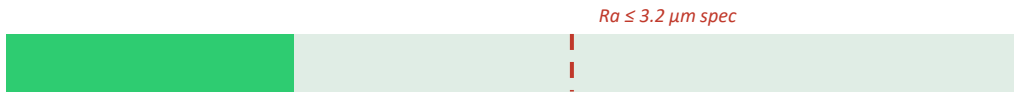

Predicted: 1.627 µm

+1.573 µm margin · Spec: ≤ 3.2 µm

✅ **Pass**

vs

Scenario B — 12 mm Steel · Under-Powered

Material	**Carbon Steel · 12 mm**
Laser Power	**2,500 W ← too low**
Cut Speed	**35 mm/s ← too fast**
Gas Flow	**5 L/min ← insufficient**
O_2 Content	**25%**

Predicted: 4.174 µm

−0.974 µm over spec · Spec: ≤ 3.2 µm

❌ **Fail**

Scenario C (corrected recipe): 3,800W · 18 mm/s · 10 L/min → Predicted 2.935 µm · ✅ Pass (+0.265 µm margin) · Recovery: 1.24 µm

The Ridge Path

All 8 features. All retained. All shrinking toward zero — but none ever reaching it.

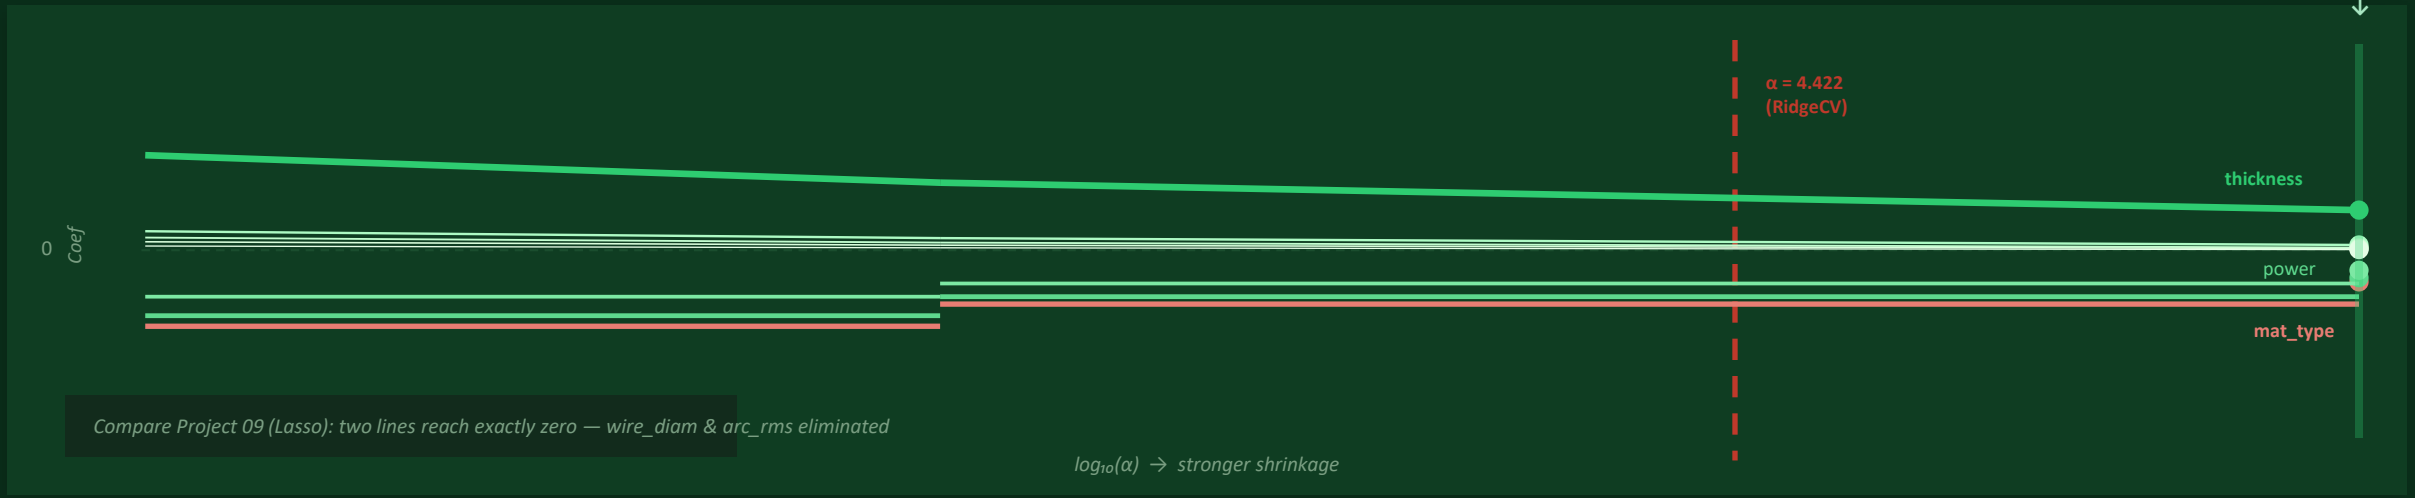

0.722

R^2 — Ridge = OLS

Identical accuracy. Ridge wins on coefficient stability — not predictive power. That is the correct trade.

7.6

Max VIF → Ridge justified

Variance Inflation Factor of 7.6 on laser_power quantifies the collinearity that makes OLS unreliable.

8/8

Features retained

Unlike Lasso, Ridge keeps all 8 variables. All have physical justification. None are discarded.

Lasso zeroes. Ridge steadies. Both are right — for different problems.

When your process variables move together because your operators are doing their job correctly, you don't need elimination. You need a model that can work with the correlation — and be trusted.

Clone it. Challenge it. Take it to your cutting cell.

Luis Lozano | Operational Excellence Manager · MBB · Machine Learning | GitHub: LozanoLsa